

Why all this?

Differential privacy is given w.r.t. datasets that "differ in only one entry", but the mechanisms we have for making a function differentially private measure the noise they add in terms of a sensitivity that is given w.r.t. e.g. the L_2 -distance. Hence we need to check sensitivity of functions w.r.t. different input- and output metrics.

Metric Spaces

(M, d) where M is a set and $d : M \times M \rightarrow \mathbb{R}$ s.t.

- $d(x, y) = 0 \Leftrightarrow x = y$
- $d(x, y) = d(y, x)$
- $d(x, z) \leq d(x, y) + d(y, z)$

Sensitivity

for metric spaces M, N a map $f : M \rightarrow N$ is s -sensitive if for all $x, y \in M$

$$d_N(f(x), f(y)) \leq s \cdot d_M(x, y)$$

"If the input is at most 1 apart, the output is at most s apart."

Gaussian Mechanism

Let \mathcal{D} be some space equipped with the discrete metric (e.g. (\mathbb{D}, L_∞)). Given a function $f : \mathcal{D} \rightarrow \mathbb{R}^n$ that is s -sensitive in L_2 norm, for every $\delta \in (0, 1)$ and $\epsilon \in (0, 1)$ the gaussian mechanism

$$\mathcal{M}_{\text{Gauss}}(f, \epsilon, \delta)(x) = f(x) + \mathcal{N}^n \left(\mu = 0, \sigma^2 = \frac{2 \ln(1.25/\delta) \cdot s^2}{\epsilon^2} \right)$$

yields an (ϵ, δ) -differentially private function.

Metrics on Numbers

- On \mathbb{R} the metric is standard:

$$d_{\mathbb{R}}(x, y) = |x - y|$$

- On \mathbb{D} we just check if the numbers are equal:

$$d_{\mathbb{D}}(x, y) = (x == y ? 0 : 1)$$

Metrics on Vectors

- Vectors over \mathbb{R} :

$$d_{L1, \mathbb{R}}(v, w) = \sum_i d_{\mathbb{R}}(v_i, w_i)$$

$$d_{L2, \mathbb{R}}(v, w) = \sqrt{\sum_i d_{\mathbb{R}}(v_i, w_i)^2}$$

$$d_{L\infty, \mathbb{R}}(v, w) = \max_i d_{\mathbb{R}}(v_i, w_i)$$

- Vectors over \mathbb{D} :

$$d_{L1, \mathbb{D}}(v, w) = \sum_i d_{\mathbb{D}}(v_i, w_i) = \text{number of entries that differ}$$

$$d_{L2,\mathbb{D}}(v, w) = \sqrt{\sum_i d_{\mathbb{D}}(v_i, w_i)^2} = \sqrt{d_{L1,\mathbb{D}}(v, w)}$$

$$d_{L\infty,\mathbb{D}}(v, w) = \max_i d_{\mathbb{D}}(v_i, w_i) = "0 \text{ if } v = w, 1 \text{ otherwise}"$$

- some facts

– for vectors $v \neq w$ we have

$$1 = d_{L\infty,\mathbb{D}}(v, w) \leq d_{L2,\mathbb{D}}(v, w) \leq d_{L1,\mathbb{D}}(v, w)$$

- so if a vector-valued function is s -sensitive for a fixed input metric and output metric $d_{L1,\mathbb{D}}$, it is also s -sensitive under output metric $d_{L\infty,\mathbb{D}}$
- all functions from $(*, \mathbb{D})$ -vectors to $(L\infty, \mathbb{D})$ -vectors are 1-sensitive because for $v \neq w$ it is $d_{L\infty,\mathbb{D}}(v, w) = 1$ and $d_{L*,\mathbb{D}}(v, w) \geq 1$ always
- Clipping $(L\infty, \mathbb{D})$ vectors is such a function and hence 1-sensitive. I don't think clipping vectors in general is 1-sensitive (even if it says so in the paper), because e.g. $d_{L1,\mathbb{D}}$ can become larger:

$$d_{L1,\mathbb{D}}([1, 1], [1, 0]) = 1$$

but for the clipped vectors

$$d_{L1,\mathbb{D}}(\text{clip}^{L1}([1, 1]), \text{clip}^{L1}([1, 0])) = d_{L1,\mathbb{D}}([0.5, 0.5], [1, 0]) = 2$$

Matrix Type

The duet matrix type has the following parameters:

$$\mathbb{M}_l^c \tau[i, j]$$

is the type of matrices where

- the matrix has i rows and j columns
- all rows have $d_{c,\mathbb{R}}(r, 0) \leq 1$ (note that this is the \mathbb{R} norm no matter what τ is. This differs from what is said on p.44 of the paper, but it makes no sense otherwise and in their implementation it's like we think, see last item of "Implications")
- the elements are of type τ and the metric is chosen accordingly
- sensitivities of variables with this type are given w.r.t. $d_{\mathbb{M}_l^* \tau}$

Metrics over Matrices

For matrices $m, n \vdash \mathbb{M}_l^* \tau$ the metric sums over rows:

$$d_{\mathbb{M}_l^* \tau}(m, n) = \sum_j d_{l,\tau}(m_j, n_j)$$

In particular,

$$d_{\mathbb{M}_{L1}^* \mathbb{D}}(m, n) = \text{number of matrix entries that differ}$$

$$d_{\mathbb{M}_{L\infty}^* \mathbb{D}}(m, n) = \text{number of matrix rows that differ somewhere}$$

discf

The function $\text{discf} : \mathbb{R} \rightarrow \mathbb{D}$ is claimed to be 1-sensitive in the paper. Taking the numbers 0.1 and 0.2 as an example, we get

$$1 = d_{\mathbb{D}}(\text{discf}(0.1), \text{discf}(0.2)) = 10 \cdot d_{\mathbb{R}}(0.1, 0.2)$$

so using our notion of sensitivity discf must be at least 10-sensitive...

I suspect they used a different definition for sensitivity, namely a function $f : M \rightarrow N$ to be s -sensitive iff

$$\max_{d_M(x,y)=1} d_N(f(x), f(y)) = s$$

This definition is equivalent to the above one if $M = \mathbb{D}$ but not in general.

convert

We can convert $\mathbb{M}_*^l \mathbb{D}$ to $\mathbb{M}_*^l \mathbb{R}$ because all rows of the first type have l -norm ≤ 1 so for any two rows m_i, n_i with $d_{*, \mathbb{D}}(m_i, n_i) \leq 1$ we have

$$d_{l, \mathbb{R}}(m_i, n_i) \leq d_{l, \mathbb{R}}(m_i, 0) + d_{l, \mathbb{R}}(0, n_i) \leq 1 + 1$$

The inequality is met, e.g. by the vectors $[1, 0]$ and $[-1, 0]$, as they have \mathbb{D} -distance 1 and $L2$ -norm 1, but $L2$ -distance 2 from each other.

This implies conversion of the rows is 2-sensitive, but in the paper it is declared 1-sensitive. Also the paper version does not preserve clipping on the matrix, even though in their interpreter the **convert** function is simply the identity.